

## Aim: Implement linear regression in Python on a given dataset.

```
In [117... import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
```

## Simple linear Regression

```
In [118... #data_no=pd.read_csv('Automobile.csv')
#df = pd.DataFrame()
d = pd.read_csv("Automobile.csv")
```

```
In [119... d
```

Out[119...

	name	mpg	cylinders	displacement	horsepower	weight	acceleration	model
0	chevrolet chevelle malibu	18.0	8.0	307.0	130.0	3504.0	12.0	
1	buick skylark 320	15.0	8.0	350.0	165.0	3693.0	11.5	
2	plymouth satellite	18.0	8.0	318.0	150.0	3436.0	11.0	
3	amc rebel sst	16.0	8.0	304.0	150.0	3433.0	12.0	
4	ford torino	17.0	8.0	302.0	140.0	3449.0	10.5	
...	...	...	...	...	...	...	...	...
393	ford mustang gl	27.0	4.0	140.0	86.0	2790.0	15.6	
394	vw pickup	44.0	4.0	97.0	52.0	2130.0	24.6	
395	dodge rampage	32.0	4.0	135.0	84.0	2295.0	11.6	
396	ford ranger	28.0	4.0	120.0	79.0	2625.0	18.6	
397	chevy s-10	31.0	4.0	119.0	82.0	2720.0	19.4	

398 rows × 9 columns



In [120...

```
d.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   name            398 non-null   object
1   mpg             397 non-null   float64
2   cylinders       397 non-null   float64
3   displacement    393 non-null   float64
4   horsepower      391 non-null   float64
5   weight          393 non-null   float64
6   acceleration    396 non-null   float64
7   model_year     396 non-null   object
8   origin          397 non-null   object
dtypes: float64(6), object(3)
memory usage: 28.1+ KB
```

In [ ]:

```
In [121... d.isnull().sum()
```

```
Out[121... name          0
mpg            1
cylinders      1
displacement   5
horsepower     7
weight         5
acceleration   2
model_year     2
origin         1
dtype: int64
```

```
In [122... X = d.iloc[:,3:4]
Y = d.iloc[:,4:5]
```

```
In [123... X=X.fillna(192.393130)
Y=Y.fillna(104.524297)
```

```
In [124... Y.shape
```

```
Out[124... (398, 1)
```

```
In [125... X.shape
```

```
Out[125... (398, 1)
```

```
In [126... from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.4, random_
```

```
In [127... Y_test.shape
```

```
Out[127... (160, 1)
```

```
In [128... X_test
```

Out[128...

displacement	
94	440.0
32	98.0
279	98.0
178	120.0
354	100.0
...	...
23	121.0
241	146.0
1	350.0
169	232.0
39	400.0

160 rows × 1 columns

In [129...


`d.describe()`

Out[129...

	mpg	cylinders	displacement	horsepower	weight	acceleration
count	397.000000	397.000000	393.000000	391.000000	393.000000	396.000000
mean	23.538539	5.458438	192.393130	104.524297	2949.053435	15.671970
std	7.811191	1.701577	103.205814	38.525101	851.576054	3.961926
min	9.000000	3.000000	68.000000	46.000000	15.000000	0.500000
25%	17.500000	4.000000	105.000000	75.000000	2220.000000	13.775000
50%	23.000000	4.000000	146.000000	94.000000	2790.000000	15.500000
75%	29.000000	8.000000	260.000000	127.000000	3570.000000	17.225000
max	46.600000	8.000000	455.000000	230.000000	5140.000000	70.000000

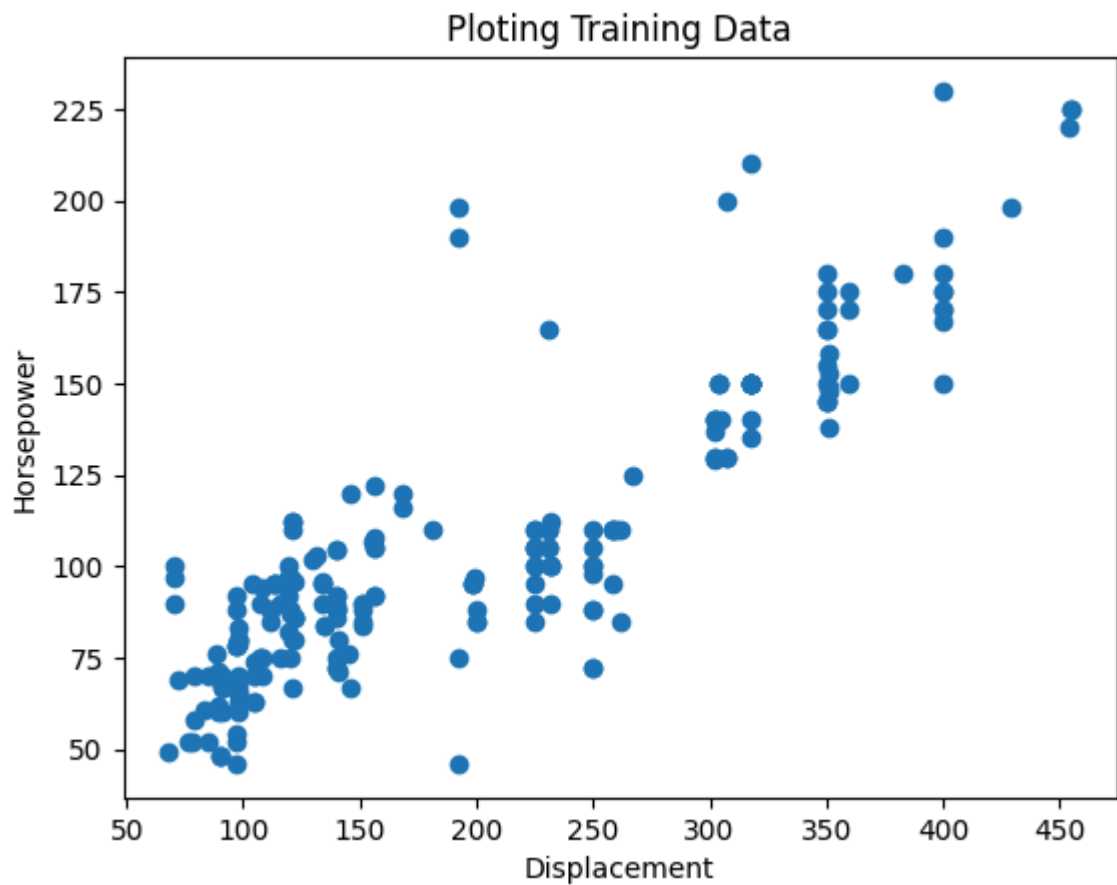
In [130...

```
plt.scatter(X,Y)
# naming the x axis
plt.xlabel('Displacement')
# naming the y axis
plt.ylabel('Horsepower')
# giving a title to graph
plt.title('All Data')
# function to show the plot
plt.show()
#plotting X & Y data on graph
```




In [131...

```
# naming the x axis
plt.xlabel('Displacement')
# naming the y axis
plt.ylabel('Horsepower')
# giving a title to graph
plt.title('Plotting Training Data')
#plotting X & Y data on graph
plt.scatter(X_train,Y_train)
plt.show()
#plotting training data on graph
```



```
In [132... #plotting Testing data on graph

# naming the x axis
plt.xlabel('Displacement')
# naming the y axis
plt.ylabel('Horsepower')
# giving a title to graph
plt.title('Ploting Testing Data')
#plotting X & Y data on graph
plt.scatter(X_test,Y_test)
plt.show()
```



```
In [133... from sklearn import linear_model
lr = linear_model.LinearRegression()
lr.fit(X_train,Y_train)

intercept=lr.intercept_
intercept
```

```
Out[133... array([40.45473373])
```

```
In [134... coef=lr.coef_
coef
```

```
Out[134... array([[0.33128207]])
```

```
In [135... Y_pred = lr.predict(X_test)
```


```
In [136... # Plotting the scatter plot
plt.scatter(X_test, Y_test, color='blue', label='Testing data')

# Plotting the regression line with predicted values
plt.plot(X_test, Y_pred, color='red', linewidth=2, label='Predicted Line')

plt.xlabel('Displacement')
plt.ylabel('Horsepower')
plt.title('Plotting Testing Data and Predicted Data')

plt.legend()

plt.show()
```



```
In [137...] lr.score(X_train, Y_train)
```

```
Out[137...] 0.7579028942658471
```

```
In [138...] Y_pred=lr.predict(X_test)
from sklearn.metrics import mean_absolute_error, mean_squared_error, root_mean_squ
mean_absolute_error(Y_test, Y_pred)
```

```
Out[138...] 13.070989131325026
```

```
In [139...] mean_squared_error(Y_test, Y_pred)
```

```
Out[139...] 367.74996277773056
```

```
In [140...] root_mean_squared_error(Y_test, Y_pred)
```

```
Out[140...] 19.176807940262908
```

```
In [141...] r2=r2_score(Y_test, Y_pred)
r2
```

```
Out[141...] 0.7345525423105541
```

```
In [142...] n=398
k=1
adj_r2score = 1 - (((1-r2)*(10-n))/(10-n-k))
print(adj_r2score)
```

```
0.7352349265205527
```



# Mutivalue linear Regression

```
In [143... #data_no=pd.read_csv('Automobile.csv')
#df = pd.DataFrame()
M2 = pd.read_csv("Automobile.csv")
```

```
In [144... M2.isnull().sum()
```

```
Out[144... name          0
mpg            1
cylinders      1
displacement   5
horsepower     7
weight         5
acceleration   2
model_year     2
origin         1
dtype: int64
```

```
In [145... d.describe()
```

```
Out[145...
```

	mpg	cylinders	displacement	horsepower	weight	acceleration
<b>count</b>	397.000000	397.000000	393.000000	391.000000	393.000000	396.000000
<b>mean</b>	23.538539	5.458438	192.393130	104.524297	2949.053435	15.671970
<b>std</b>	7.811191	1.701577	103.205814	38.525101	851.576054	3.961926
<b>min</b>	9.000000	3.000000	68.000000	46.000000	15.000000	0.500000
<b>25%</b>	17.500000	4.000000	105.000000	75.000000	2220.000000	13.775000
<b>50%</b>	23.000000	4.000000	146.000000	94.000000	2790.000000	15.500000
<b>75%</b>	29.000000	8.000000	260.000000	127.000000	3570.000000	17.225000
<b>max</b>	46.600000	8.000000	455.000000	230.000000	5140.000000	70.000000

```
In [146... X = M2.iloc[:,2:4]
Y = M2.iloc[:,4:5]
```

```
In [147... X.isnull().sum()
```

```
Out[147... cylinders      1
displacement     5
dtype: int64
```

```
In [148... Y.isnull().sum()
```

```
Out[148... horsepower    7
dtype: int64
```

```
In [149... # Replace null values in the 'cylinders' column with a specific value
X['cylinders'] = X['cylinders'].fillna(value=5.458438)

# Replace null values in the 'displacement' column with a different value
X['displacement'] = X['displacement'].fillna(value=192.393130)
```

```
Y=Y.fillna(104.524297)
```

```
In [150... X.isnull().sum()
```

```
Out[150... cylinders      0  
displacement    0  
dtype: int64
```

```
In [151... X.shape
```

```
Out[151... (398, 2)
```

```
In [152... Y.shape
```

```
Out[152... (398, 1)
```

```
In [153... from sklearn.model_selection import train_test_split  
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.4, random_
```

```
In [154... Y_test.shape
```

```
Out[154... (160, 1)
```

```
In [155... X_test.shape
```

```
Out[155... (160, 2)
```

```
In [156... from sklearn import linear_model  
lr = linear_model.LinearRegression()  
lr.fit(X_train,Y_train)  
  
intercept=lr.intercept_  
intercept
```

```
Out[156... array([28.39757549])
```

```
In [157... coef_cylinders, coef_displacement = lr.coef_[0]  
print(coef_cylinders, coef_displacement)
```

```
4.982040075188673 0.25333950531970395
```

```
In [158... coef=lr.coef_  
coef
```

```
Out[158... array([[4.98204008, 0.25333951]])
```

```
In [159... Y_pred = lr.predict(X_test)
```


```
In [160... from mpl_toolkits.mplot3d import Axes3D  
import matplotlib.pyplot as plt  
import numpy as np
```

```
In [161... # Creating a meshgrid for the plane  
x_values = np.linspace(min(X["cylinders"]), max(X["cylinders"]), 10)  
y_values = np.linspace(min(X["displacement"]), max(X["displacement"]), 10)  
x_mesh, y_mesh = np.meshgrid(x_values, y_values)
```

```
In [162... # Calculating z values for the plane using the plane equation
z_mesh = coef_cylinders * x_mesh + coef_displacement * y_mesh + intercept
```

```
In [163... # Plotting the scatter plot
fig = plt.figure()
ax = fig.add_subplot(projection='3d')
ax.scatter(X["cylinders"], X["displacement"], Y)
# Plotting the best fit plane
ax.plot_surface(x_mesh, y_mesh, z_mesh, alpha = 0.3, color='green')
ax.set_xlabel('Cylinders')
ax.set_ylabel('Displacement')
ax.set_zlabel('Horsepower')
ax.set_title('3D Scatter Plot with Best Fit Plane')
# ax.legend()
plt.show()
```

### 3D Scatter Plot with Best Fit Plane



```
In [164... Y_pred=lr.predict(X_test)
from sklearn.metrics import mean_absolute_error, mean_squared_error, root_mean_squ
mean_absolute_error(Y_test, Y_pred)
```

```
Out[164... 13.34665612540907
```

```
In [165... mean_squared_error(Y_test, Y_pred)
```

```
Out[165... 363.8637958562596
```

```
In [166... root_mean_squared_error(Y_test, Y_pred)
```

```
Out[166... 19.075214175894846
```

```
In [167... r2=r2_score(Y_test, Y_pred)
r2
```

Out[167... 0.7373576360804337

In [168... Y\_test.shape

Out[168... (160, 1)

```
In [169... n=160
k=2
adj_r2score = 1 - (((1-r2)*(10-n))/(10-n-k))
print(adj_r2score)
```

0.7408134566583227

## More than 2 Mutivalue linear Regression

```
In [170... #data_no=pd.read_csv('Automobile.csv')
#M4 = pd.DataFrame()
M4 = pd.read_csv("Automobile.csv")
```

In [171... M4

Out[171...

	name	mpg	cylinders	displacement	horsepower	weight	acceleration	model
<b>0</b>	chevrolet chevelle malibu	18.0	8.0	307.0	130.0	3504.0	12.0	
<b>1</b>	buick skylark 320	15.0	8.0	350.0	165.0	3693.0	11.5	
<b>2</b>	plymouth satellite	18.0	8.0	318.0	150.0	3436.0	11.0	
<b>3</b>	amc rebel sst	16.0	8.0	304.0	150.0	3433.0	12.0	
<b>4</b>	ford torino	17.0	8.0	302.0	140.0	3449.0	10.5	
...	...	...	...	...	...	...	...	...
<b>393</b>	ford mustang gl	27.0	4.0	140.0	86.0	2790.0	15.6	
<b>394</b>	vw pickup	44.0	4.0	97.0	52.0	2130.0	24.6	
<b>395</b>	dodge rampage	32.0	4.0	135.0	84.0	2295.0	11.6	
<b>396</b>	ford ranger	28.0	4.0	120.0	79.0	2625.0	18.6	
<b>397</b>	chevy s- 10	31.0	4.0	119.0	82.0	2720.0	19.4	

398 rows × 9 columns



In [172...

M4.isnull().sum()

Out[172...

```

name          0
mpg           1
cylinders     1
displacement  5
horsepower    7
weight        5
acceleration  2
model_year    2
origin        1
dtype: int64

```

In [173...

M4.describe()

Out[173...

	mpg	cylinders	displacement	horsepower	weight	acceleration
<b>count</b>	397.000000	397.000000	393.000000	391.000000	393.000000	396.000000
<b>mean</b>	23.538539	5.458438	192.393130	104.524297	2949.053435	15.671970
<b>std</b>	7.811191	1.701577	103.205814	38.525101	851.576054	3.961926
<b>min</b>	9.000000	3.000000	68.000000	46.000000	15.000000	0.500000
<b>25%</b>	17.500000	4.000000	105.000000	75.000000	2220.000000	13.775000
<b>50%</b>	23.000000	4.000000	146.000000	94.000000	2790.000000	15.500000
<b>75%</b>	29.000000	8.000000	260.000000	127.000000	3570.000000	17.225000
<b>max</b>	46.600000	8.000000	455.000000	230.000000	5140.000000	70.000000

In [174...

M4.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   name            398 non-null    object
1   mpg             397 non-null    float64
2   cylinders       397 non-null    float64
3   displacement    393 non-null    float64
4   horsepower      391 non-null    float64
5   weight          393 non-null    float64
6   acceleration    396 non-null    float64
7   model_year      396 non-null    object
8   origin          397 non-null    object
dtypes: float64(6), object(3)
memory usage: 28.1+ KB
```

In [175...

```
from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
```

In [176...

```
# new_Crash_data=pd.DataFrame()
M4['name'] = encoder.fit_transform(M4['name'])
# new_Crash_data=pd.DataFrame()
M4['model_year'] = encoder.fit_transform(M4['model_year'])
# new_Crash_data=pd.DataFrame()
M4['origin'] = encoder.fit_transform(M4['origin'])
```

In [177...

M4

Out[177...

	name	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year
<b>0</b>	49	18.0	8.0	307.0	130.0	3504.0	12.0	
<b>1</b>	36	15.0	8.0	350.0	165.0	3693.0	11.5	
<b>2</b>	231	18.0	8.0	318.0	150.0	3436.0	11.0	
<b>3</b>	14	16.0	8.0	304.0	150.0	3433.0	12.0	
<b>4</b>	161	17.0	8.0	302.0	140.0	3449.0	10.5	
...	...	...	...	...	...	...	...	...
<b>393</b>	153	27.0	4.0	140.0	86.0	2790.0	15.6	
<b>394</b>	301	44.0	4.0	97.0	52.0	2130.0	24.6	
<b>395</b>	119	32.0	4.0	135.0	84.0	2295.0	11.6	
<b>396</b>	159	28.0	4.0	120.0	79.0	2625.0	18.6	
<b>397</b>	69	31.0	4.0	119.0	82.0	2720.0	19.4	

398 rows × 9 columns



In [178...

```
import pandas as pd

X = pd.concat([M4.iloc[:,0:4], M4.iloc[:,5:9]], axis=1)
```

In [179...

```
M4.isnull().sum()
```

Out[179...

```
name          0
mpg           1
cylinders     1
displacement  5
horsepower    7
weight        5
acceleration  2
model_year    0
origin        0
dtype: int64
```

In [180...

```
X
```

Out[180...

	name	mpg	cylinders	displacement	weight	acceleration	model_year	origin
0	49	18.0	8.0	307.0	3504.0	12.0	0	2
1	36	15.0	8.0	350.0	3693.0	11.5	0	2
2	231	18.0	8.0	318.0	3436.0	11.0	0	2
3	14	16.0	8.0	304.0	3433.0	12.0	0	2
4	161	17.0	8.0	302.0	3449.0	10.5	0	2
...	...	...	...	...	...	...	...	...
393	153	27.0	4.0	140.0	2790.0	15.6	12	2
394	301	44.0	4.0	97.0	2130.0	24.6	12	0
395	119	32.0	4.0	135.0	2295.0	11.6	12	2
396	159	28.0	4.0	120.0	2625.0	18.6	12	2
397	69	31.0	4.0	119.0	2720.0	19.4	12	2

398 rows × 8 columns

In [181...

X.describe()

Out[181...

	name	mpg	cylinders	displacement	weight	acceleration	mc
count	398.000000	397.000000	397.000000	393.000000	393.000000	396.000000	39
mean	148.550251	23.538539	5.458438	192.393130	2949.053435	15.671970	
std	89.495880	7.811191	1.701577	103.205814	851.576054	3.961926	
min	0.000000	9.000000	3.000000	68.000000	15.000000	0.500000	
25%	65.250000	17.500000	4.000000	105.000000	2220.000000	13.775000	
50%	150.000000	23.000000	4.000000	146.000000	2790.000000	15.500000	
75%	225.750000	29.000000	8.000000	260.000000	3570.000000	17.225000	
max	304.000000	46.600000	8.000000	455.000000	5140.000000	70.000000	1



In [182...

```

X['mpg'] = X['mpg'].fillna(value=23.538539)
# Replace null values in the 'cylinders' column with a specific value
X['cylinders'] = X['cylinders'].fillna(value=5.458438)

# Replace null values in the 'displacement' column with a different value
X['displacement'] = X['displacement'].fillna(value=192.393130)

X['weight'] = X['weight'].fillna(value=192.393130)

X['acceleration'] = X['acceleration'].fillna(value=192.393130)

Y=Y.fillna(104.524297)

```



```

In [183... from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.4, random_

In [184... from sklearn import linear_model
lr = linear_model.LinearRegression()
lr.fit(X_train,Y_train)

intercept=lr.intercept_
intercept

Out[184... array([64.13840453])

In [185... coef=lr.coef_
coef

Out[185... array([[ -0.01040342, -0.91356548,  2.70751266,  0.18423884,  0.00718577,
        -0.10747737, -0.361584   , -2.82215043]])

In [186... coef_name, coef_mpg, coef_cylinders, coef_displacement, coef_weight, coef_accele
print(coef_mpg, coef_cylinders, coef_displacement)

-0.9135654830247643 2.707512656282228 0.18423883622165788

In [187... Y_pred = lr.predict(X_test)

In [188... Y_pred=lr.predict(X_test)
from sklearn.metrics import mean_absolute_error,mean_squared_error,root_mean_squ
mean_absolute_error(Y_test, Y_pred)

Out[188... 13.295338700617416

In [189... mean_squared_error(Y_test, Y_pred)

Out[189... 341.332185874115

In [190... root_mean_squared_error(Y_test, Y_pred)

Out[190... 18.475177560015897

In [191... r2=r2_score(Y_test, Y_pred)
r2

Out[191... 0.7536212912613462

In [192... n=160
k=3
adj_r2score = 1 - (((1-r2)*(10-n))/(10-n-k))
print(adj_r2score)

0.7584522463346531

```

## Removing origin column

```

In [193... # column_to_drop = "origin"
X = X.drop('origin', axis=1)

```

In [194...

X

Out[194...

	name	mpg	cylinders	displacement	weight	acceleration	model_year
0	49	18.0	8.0	307.0	3504.0	12.0	0
1	36	15.0	8.0	350.0	3693.0	11.5	0
2	231	18.0	8.0	318.0	3436.0	11.0	0
3	14	16.0	8.0	304.0	3433.0	12.0	0
4	161	17.0	8.0	302.0	3449.0	10.5	0
...	...	...	...	...	...	...	...
393	153	27.0	4.0	140.0	2790.0	15.6	12
394	301	44.0	4.0	97.0	2130.0	24.6	12
395	119	32.0	4.0	135.0	2295.0	11.6	12
396	159	28.0	4.0	120.0	2625.0	18.6	12
397	69	31.0	4.0	119.0	2720.0	19.4	12

398 rows × 7 columns

In [195...

```
X['mpg'] = X['mpg'].fillna(value=23.538539)
# Replace null values in the 'cylinders' column with a specific value
X['cylinders'] = X['cylinders'].fillna(value=5.458438)

# Replace null values in the 'displacement' column with a different value
X['displacement'] = X['displacement'].fillna(value=192.393130)

X['weight'] = X['weight'].fillna(value=192.393130)

X['acceleration'] = X['acceleration'].fillna(value=192.393130)

Y=Y.fillna(104.524297)
```

In [196...

```
from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.4, random_
```

In [197...

```
from sklearn import linear_model
lr = linear_model.LinearRegression()
lr.fit(X_train,Y_train)

intercept=lr.intercept_
intercept
```

Out[197...

array([58.6808647])

In [198...

```
coef=lr.coef_
coef
```

Out[198...

```
array([[ -0.00420697, -0.84155477,  2.78998424,  0.17328983,  0.0075213 ,
        -0.10418774, -0.48523644]])
```

```
In [199... coef_name, coef_mpg, coef_cylinders, coef_displacement, coef_weight, coef_accele
print(coef_mpg, coef_cylinders, coef_displacement)
```

```
-0.8415547657354356 2.7899842356934568 0.17328982724531156
```

```
In [200... Y_pred = lr.predict(X_test)
```

```
In [201... Y_pred=lr.predict(X_test)
from sklearn.metrics import mean_absolute_error, mean_squared_error, root_mean_squ
mean_absolute_error(Y_test, Y_pred)
```

```
Out[201... 13.35629564001081
```

```
In [202... r2=r2_score(Y_test, Y_pred)
r2
```

```
Out[202... 0.7464578917754212
```

## Removing Acceleration Column

```
In [203... # column_to_drop = "origin"
X = X.drop('acceleration', axis=1)
```

```
In [204... from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.4, random_
```

```
In [205... from sklearn import linear_model
lr = linear_model.LinearRegression()
lr.fit(X_train, Y_train)
```

```
Out[205... LinearRegression ⓘ ?
LinearRegression()
```

```
In [208... Y_pred = lr.predict(X_test)
```

```
In [209... Y_pred=lr.predict(X_test)
from sklearn.metrics import mean_absolute_error, mean_squared_error, root_mean_squ
mean_absolute_error(Y_test, Y_pred)
r2=r2_score(Y_test, Y_pred)
r2
```

```
Out[209... 0.7420533440568693
```

```
In [ ]:
```